

# Operating Systems

## CS-C3140, Lecture 3

Alexandru Paler

# Fault Tolerance

- Refers to the ability of a system or component to continue normal operation despite the presence of hardware or software faults
  - involves some degree of redundancy
  - Intended to increase the reliability of a system
  - comes with a cost in financial terms or performance
- The extent adoption of fault tolerance measures must be determined by how critical the resource is

## Basic Concepts and Taxonomy of Dependable and Secure Computing

Algirdas Avizienis, *Fellow, IEEE*, Jean-Claude Laprie, Brian Randell, and Carl Landwehr

**Abstract**—This paper gives the main definitions relating to dependability, a generic concept including as special case such attributes as reliability, availability, safety, integrity, maintainability, etc. Security brings in concerns for confidentiality, in addition to availability and integrity. Basic definitions are given first. They are then commented upon, and supplemented by additional definitions, which address the threats to dependability and security (faults, errors, failures), their attributes, and the means for their achievement (fault prevention, fault tolerance, fault removal, fault forecasting). The aim is to explicate a set of general concepts, of relevance across a wide range of situations and, therefore, helping communication and cooperation among a number of scientific and technical communities, including ones that are concentrating on particular types of system, of system failures, or of causes of system failures.

**Index Terms**—Dependability, security, trust, faults, errors, failures, vulnerabilities, attacks, fault tolerance, fault removal, fault forecasting.

### 1 INTRODUCTION

THIS paper aims to give precise definitions characterizing the various concepts that come into play when addressing the dependability and security of computing and communication systems. Clarifying these concepts is surprisingly difficult when we discuss systems in which there are uncertainties about system boundaries. Furthermore, the very complexity of systems (and their specification) is often a major problem, the determination of possible causes or consequences of failure can be a very subtle process, and there are (fallible) provisions for preventing faults from causing failures.

Dependability is first introduced as a global concept that subsumes the usual attributes of reliability, availability, safety, integrity, maintainability, etc. The consideration of security brings in concerns for confidentiality, in addition to availability and integrity. The basic definitions are then commented upon and supplemented by additional definitions. **Boldface** characters are used when a term is defined, while *italic* characters are an invitation to focus the reader's attention.

This paper can be seen as an attempt to document a minimum consensus on concepts within various specialties in order to facilitate fruitful technical interactions; in addition, we hope that it will be suitable 1) for use by

the concepts: words are only of interest because they unequivocally label concepts and enable ideas and viewpoints to be shared. An important issue, for which we believe a consensus has not yet emerged, concerns the measures of dependability and security; this issue will necessitate further elaboration before being documented consistently with the other aspects of the taxonomy that is presented here.

The paper has no pretension of documenting the state-of-the-art. Thus, together with the focus on concepts, we do not address implementation issues such as can be found in standards, for example, in [30] for safety or [32] for security.

The dependability and security communities have followed distinct, but convergent paths: 1) dependability has realized that restriction to nonmalicious faults was addressing only a part of the problem, 2) security has realized that the main focus that was put in the past on confidentiality needed to be augmented with concerns for integrity and for availability (they have been always present in the definitions, but did not receive as much attention as confidentiality). The paper aims to bring together the common strands of dependability and security although, for reasons of space limitation, confidentiality is not given the attention it deserves.

# From Failure to Disaster

- Computers are components in larger technical or societal systems
- Failure detection and manual back-up system can prevent disaster
  - Used routinely in safety-critical systems
  - Manual control/override in jetliners
  - Ground-based control for spacecraft
  - Manual bypass in nuclear reactors
- Catastrophic
  - Serious consequences
- Major
  - Incorrect operation
  - Possibly recoverable
- Minor
  - Inconvenience
- Not noticed



## Space Shuttle Challenger Disaster


- Second shuttle after Columbia, first mission on April 4, 1983
- Exploded 73 seconds after launch on its 10th mission (January 28, 1986)
- Seven crew members killed
- Temperatures dipped below freezing on launch day
- Engineers determined that the cold temperatures caused a loss of flexibility in the O-rings that decreased their ability to seal the field joints, which allowed hot gas and soot to flow past the primary O-rings

# Fault-Tolerance Steps

- Fault Detection
  - Determining that a fault has occurred
- Diagnosis
  - Determining what caused the fault, or exactly which subsystem or component is faulty
- Containment
  - Prevents the propagation of faults from their origin at one point in a system to a point where it can have an effect on the service to the user
- Masking
  - Insuring that only correct values get passed to the system boundary in spite of a failed component.
- Compensation
  - If a fault occurs and is confined to a subsystem, it may be necessary for the system to provide a response to compensate for output of the faulty subsystem.
- Repair
  - The process in which faults are removed from a system. In well-designed fault tolerant systems, faults are contained before they propagate to the extent that the delivery of system service is affected.
  - This leaves a portion of the system unusable because of residual faults.
  - If subsequent faults occur, the system may be unable to cope because of this loss of resources, unless these resources are reclaimed through a recovery process which insures that no faults remain in system resources or in the system state


# Fundamental Concepts

- Reliability  $R(t)$ 
  - continuity of correct service
  - The probability of its correct operation up to time  $t$  given that the system was operating correctly at time  $t=0$
- Mean time to failure (MTTF)
  - Mean time to repair (MTTR) is the average time it takes to repair or replace a faulty element
- Availability
  - readiness for correct service
  - The fraction of time the system is available to service users' requests
- Safety: absence of catastrophic consequences on the user(s) and the environment
- Integrity: absence of improper system alterations.
- Maintainability: ability to undergo modifications and repairs



[https://en.wikipedia.org/wiki/Bathtub\\_curve](https://en.wikipedia.org/wiki/Bathtub_curve)

# Example and Classes



**Figure 2.13** System Operational States

Class	Availability	Annual Downtime
Continuous	1.0	0
Fault Tolerant	0.99999	5 minutes
Fault Resilient	0.9999	53 minutes
High Availability	0.999	8.3 hours
Normal Availability	0.99 - 0.995	44-87 hours

# Fault Categories

- Permanent
  - A fault that, after it occurs, is always present
  - The fault persists until the faulty component is replaced or repaired
- Temporary
  - A fault that is not present all the time for all operating conditions
  - Can be classified as
    - Transient – a fault that occurs only once
    - Intermittent – a fault that occurs at multiple, unpredictable times

## Basic Concepts and Taxonomy of Dependable and Secure Computing

Algirdas Avizienis, *Fellow, IEEE*, Jean-Claude Laprie, Brian Randell, and Carl Landwehr

**Abstract**—This paper gives the main definitions relating to dependability, a generic concept including as special case such attributes as reliability, availability, safety, integrity, maintainability, etc. Security brings in concerns for confidentiality, in addition to availability and integrity. Basic definitions are given first. They are then commented upon, and supplemented by additional definitions, which address the threats to dependability and security (faults, errors, failures), their attributes, and the means for their achievement (fault prevention, fault tolerance, fault removal, fault forecasting). The aim is to explicate a set of general concepts, of relevance across a wide range of situations and, therefore, helping communication and cooperation among a number of scientific and technical communities, including ones that are concentrating on particular types of system, of system failures, or of causes of system failures.

**Index Terms**—Dependability, security, trust, faults, errors, failures, vulnerabilities, attacks, fault tolerance, fault removal, fault forecasting.

### 1 INTRODUCTION

THIS paper aims to give precise definitions characterizing the various concepts that come into play when addressing the dependability and security of computing and communication systems. Clarifying these concepts is surprisingly difficult when we discuss systems in which there are uncertainties about system boundaries. Furthermore, the very complexity of systems (and their specification) is often a major problem, the determination of possible causes or consequences of failure can be a very subtle process, and there are (fallible) provisions for preventing faults from causing failures.

Dependability is first introduced as a global concept that subsumes the usual attributes of reliability, availability, safety, integrity, maintainability, etc. The consideration of security brings in concerns for confidentiality, in addition to availability and integrity. The basic definitions are then commented upon and supplemented by additional definitions. **Boldface** characters are used when a term is defined, while *italic* characters are an invitation to focus the reader's attention.

This paper can be seen as an attempt to document a minimum consensus on concepts within various specialties in order to facilitate fruitful technical interactions; in addition, we hope that it will be suitable 1) for use by

the concepts: words are only of interest because they unequivocally label concepts and enable ideas and viewpoints to be shared. An important issue, for which we believe a consensus has not yet emerged, concerns the measures of dependability and security; this issue will necessitate further elaboration before being documented consistently with the other aspects of the taxonomy that is presented here.

The paper has no pretension of documenting the state-of-the-art. Thus, together with the focus on concepts, we do not address implementation issues such as can be found in standards, for example, in [30] for safety or [32] for security.

The dependability and security communities have followed distinct, but convergent paths: 1) dependability has realized that restriction to nonmalicious faults was addressing only a part of the problem, 2) security has realized that the main focus that was put in the past on confidentiality needed to be augmented with concerns for integrity and for availability (they have been always present in the definitions, but did not receive as much attention as confidentiality). The paper aims to bring together the common strands of dependability and security although, for reasons of space limitation, confidentiality is not given the attention it deserves.

# Causes of Errors

- Nondeterminate program operation
  - When **programs share memory**, and their execution is interleaved by the processor, they may interfere with each other by overwriting common memory areas in unpredictable ways
  - **The order in which programs are scheduled** may affect the outcome of any particular program
- Improper synchronization
  - It is often the case that **a routine must be suspended awaiting an event elsewhere in the system**
  - Improper design of the signaling mechanism can result in loss or duplication
- Failed mutual exclusion
  - More than one user or program attempts to **make use of a shared resource at the same time**
  - There must be some sort of mutual exclusion mechanism that permits only one routine at a time to perform an update against the file
- **Deadlocks**
  - It is possible for two or more programs to be hung up waiting for each other



# Dependability and RAS (Reliability, Availability and Serviceability)

- A system should
  - be capable of detecting hardware errors, and, when possible correcting them in runtime
  - provide mechanisms to detect hardware degradation, in order to warn the system administrator to take the action of replacing a component before it causes data loss or system downtime
- **Dependability**
  - original definition: is the ability to deliver service that can justifiably be trusted
  - alternate definition: the ability to avoid service failures that are more frequent and more severe than is acceptable
- The dependence of system A on system B, thus, represents the extent to which system A's dependability is (or would be) affected by that of System B.
- The concept of dependence leads to that of trust, which can very conveniently be defined as accepted dependence.

Concept	Dependability	High Confidence	Survivability	Trustworthiness
Goal	1) ability to deliver service that can justifiably be trusted 2) ability of a system to avoid service failures that are more frequent or more severe than is acceptable	consequences of the system behavior are well understood and predictable	capability of a system to fulfill its mission in a timely manner	assurance that a system will perform as expected
Threats present	1) development faults (e.g., software flaws, hardware errata, malicious logic) 2) physical faults (e.g., production defects, physical deterioration) 3) interaction faults (e.g., physical interference, input mistakes, attacks, including viruses, worms, intrusions)	• internal and external threats • naturally occurring hazards and malicious attacks from a sophisticated and well-funded adversary	1) attacks (e.g., intrusions, probes, denials of service) 2) failures (internally generated events due to, e.g., software design errors, hardware degradation, human errors, corrupted data) 3) accidents (externally generated events such as natural disasters)	1) hostile attacks (from hackers or insiders) 2) environmental disruptions (accidental disruptions, either man-made or natural) 3) human and operator errors (e.g., software flaws, mistakes by human operators)
Reference	This paper	"Information Technology Frontiers for a New Millennium (Blue Book 2000)" (48)	"Survivable network systems" (16)	"Trust in cyberspace" (62)

Fig. 15. Dependability, high confidence, survivability, and trustworthiness.

# Methods of Redundancy - No single point of failure

- Design the system with multiple instances of critical units in such a manner that the failure of some of these units does not directly fail the entire system.
- A number of techniques can be incorporated into OS software to support fault tolerance:
  - Process isolation
  - Concurrency controls
  - Virtual machines
  - Checkpoints and rollbacks

## Spatial (physical) redundancy

Involves the use of multiple components that either perform the same function simultaneously or are configured so that one component is available as a backup in case of the failure of another component

## Temporal redundancy

Involves repeating a function or operation when an error is detected

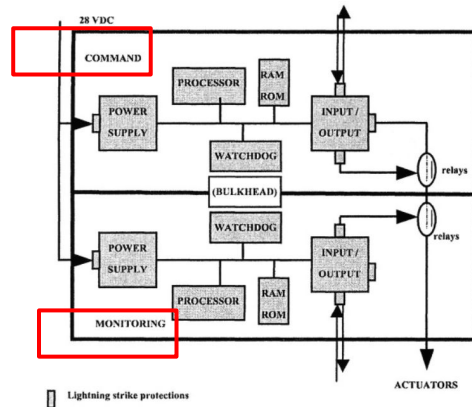
Is effective with temporary faults but not useful for permanent faults

## Information redundancy

Provides fault tolerance by replicating or coding data in such a way that bit errors can be both detected and corrected

# Example: Fly-by-wire and redundancy

- The first electrical flight control system (a.k.a. Fly-by-Wire) for a civil aircraft was designed by Aerospatiale and installed on Concorde.
- Airbus A320 entered operation in 1988
  - Other models include A340 and A380
  - Primary (P) and secondary (S) computers (different designs and suppliers)
  - Multiple redundant software modules



Two types of computers

- PRIM's (primary computers)
- SEC's (secondary computers)


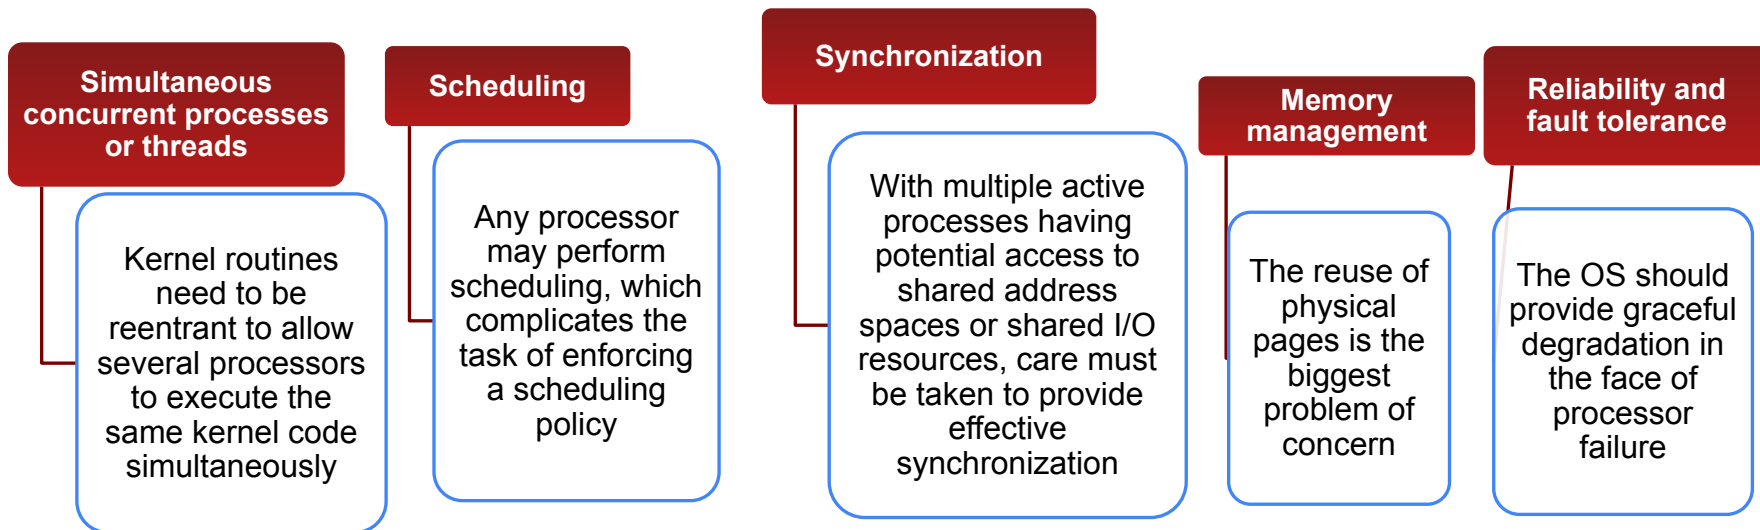


Figure 2: A340-600 system architecture

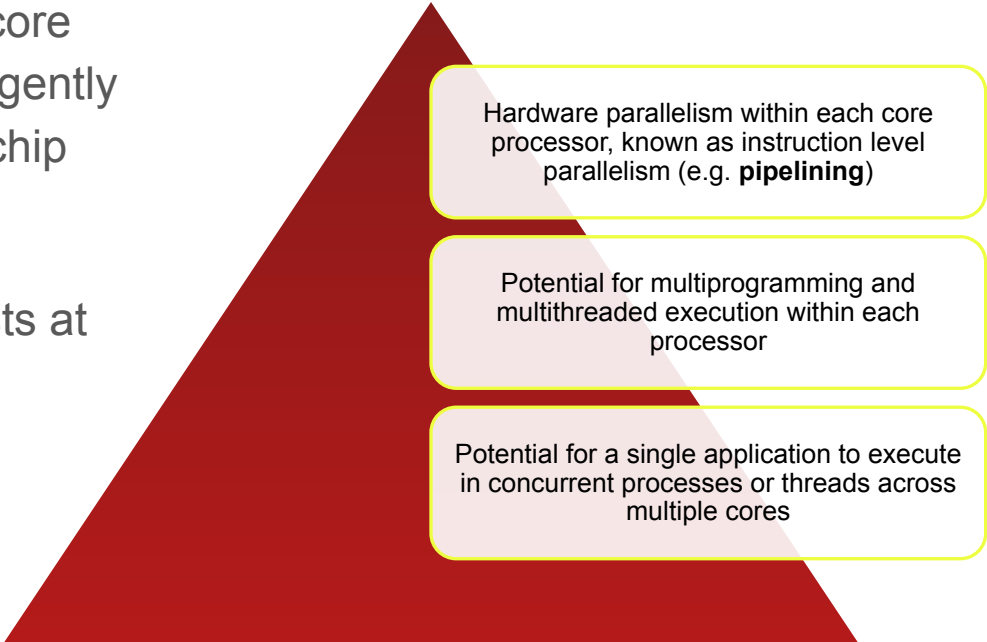
# Symmetric Multiprocessor OS Considerations

- A multiprocessor OS must provide all the functionality of a multiprogramming system plus additional features to accommodate multiple processors
- **Key design issues:**



# Parallelism and Multicore OS Considerations

- The design challenge for a many-core multicore system is to efficiently harness the multicore processing power and intelligently manage the substantial on-chip resources efficiently
- Potential for parallelism exists at three levels:




Hardware parallelism within each core processor, known as instruction level parallelism (e.g. **pipelining**)

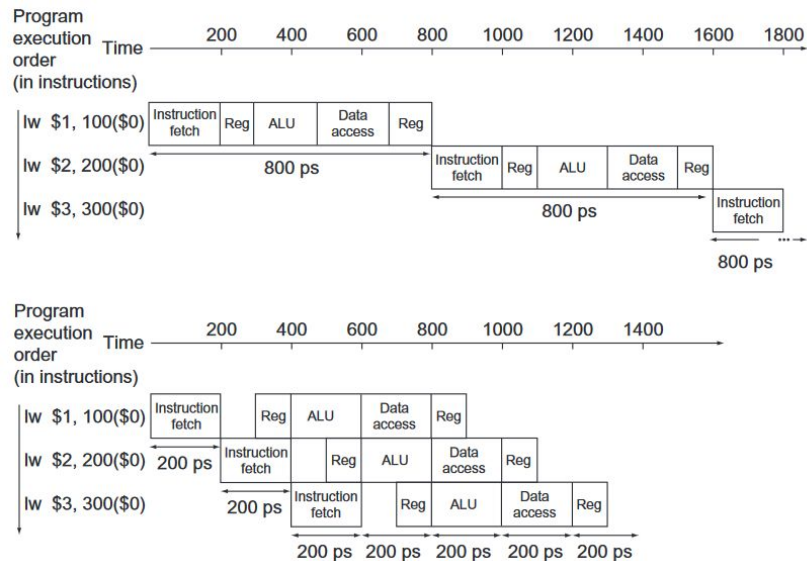
Potential for multiprogramming and multithreaded execution within each processor

Potential for a single application to execute in concurrent processes or threads across multiple cores

# Instruction Level Parallelism (ILP): Pipelining



**FIGURE 4.25 The laundry analogy for pipelining.** Ann, Brian, Cathy, and Don each have dirty clothes to be washed, dried, folded, and put away. The washer, dryer, “folder,” and “storer” each take 30 minutes for their task. Sequential laundry takes 8 hours for 4 loads of wash, while pipelined laundry takes just 3.5 hours. We show the pipeline stage of different loads over time by showing duplicate copies of the four resources on this two-dimensional time line, but we really have just one of each resource.



**FIGURE 4.27 Single-cycle, nonpipelined execution in top versus pipelined execution in bottom.** Both use the same hardware components, whose time is listed in Figure 4.26. In this case, we see a fourfold speed-up on average time between instructions, from 800 ps down to 200 ps. Compare this figure to Figure 4.25. For the laundry, we assumed all stages were equal. If the dryer were slowest, then the dryer stage would set the stage time. The pipeline stage times of a computer are also limited by the slowest resource, either the ALU operation or the memory access. We assume the write to the register file occurs in the first half of the clock cycle and the read from the register file occurs in the second half. We use this assumption throughout this chapter.


# Virtualization and Clouds

- Traditional cloud computing is about
  - IaaS (Infrastructure as a Service)
    - Virtual machines, storage, communication and balancing, etc.
  - PaaS (Platform as a Service)
    - Runtimes, servers, etc. and the related development tools
  - SaaS (Software as a Service)
    - Basically, on-demand software
- Also, cluster computing uses layers
  - Technically many such setup are based on virtualization
- Virtualization: using virtual (abstract) resources instead of the concrete resources (on which they are based)
- Allows one or more cores to be dedicated to a particular process and then leave the processor alone to devote its efforts to that process
- Multicore OS could then act as a hypervisor that makes a high-level decision to allocate cores to applications but does little in the way of resource allocation beyond that

# Virtualization - Hypervisor


The virtual machine monitor (VMM), or **hypervisor**

- runs on top of (or is incorporated into) the host OS
- supports VMs, which are emulated hardware devices
- each VM runs a separate OS
- handles each operating system's communications with the processor, the storage medium, and the network
- hands off the processor control to a virtual OS on a VM

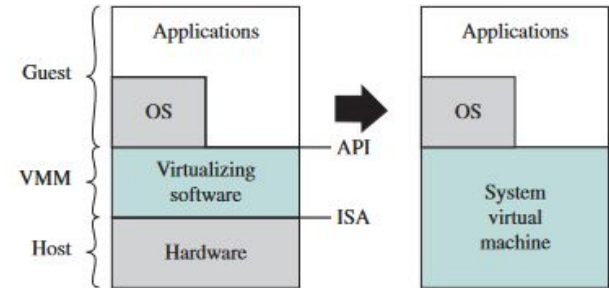


PikeOS is a combination of hypervisor and real-time operating system and is ARINC 653 standard compliant.

ARINC 653 standard defines virtualization in terms of static resources (memory, I/O) as well as processor time. It also defines an API suitable for Avionics concepts. The Airbus A380 and the Boeing A787 are examples of the IMA concept success.



(a) Process VM



(b) System VM

Figure 2.14 Process and System Virtual Machines



# Traditional UNIX Systems

- Developed at Bell Labs and became operational on a PDP-7 in 1970
- The first notable milestone was porting the UNIX system from the PDP-7 to the PDP-11
- First showed that UNIX would be an OS for all computers
- **Next milestone was rewriting UNIX in the programming language C**
- Demonstrated the advantages of using a high-level language for system code
- Was described in a technical journal for the first time in 1974
- First widely available version outside Bell Labs was Version 6 in 1976
- Version 7, released in 1978, is the ancestor of most modern UNIX systems
- Most important of the non-AT&T systems was UNIX BSD (Berkeley Software Distribution), running first on PDP and then on VAX computers

## The UNIX Time-Sharing System\*

*D. M. Ritchie and K. Thompson*

### ABSTRACT

Unix is a general-purpose, multi-user, interactive operating system for the larger Digital Equipment Corporation PDP-11 and the Interdata 8/32 computers. It offers a number of features seldom found even in larger operating systems, including

- i A hierarchical file system incorporating demountable volumes,
- ii Compatible file, device, and inter-process I/O,
- iii The ability to initiate asynchronous processes,
- iv System command language selectable on a per-user basis,
- v Over 100 subsystems including a dozen languages,
- vi High degree of portability.

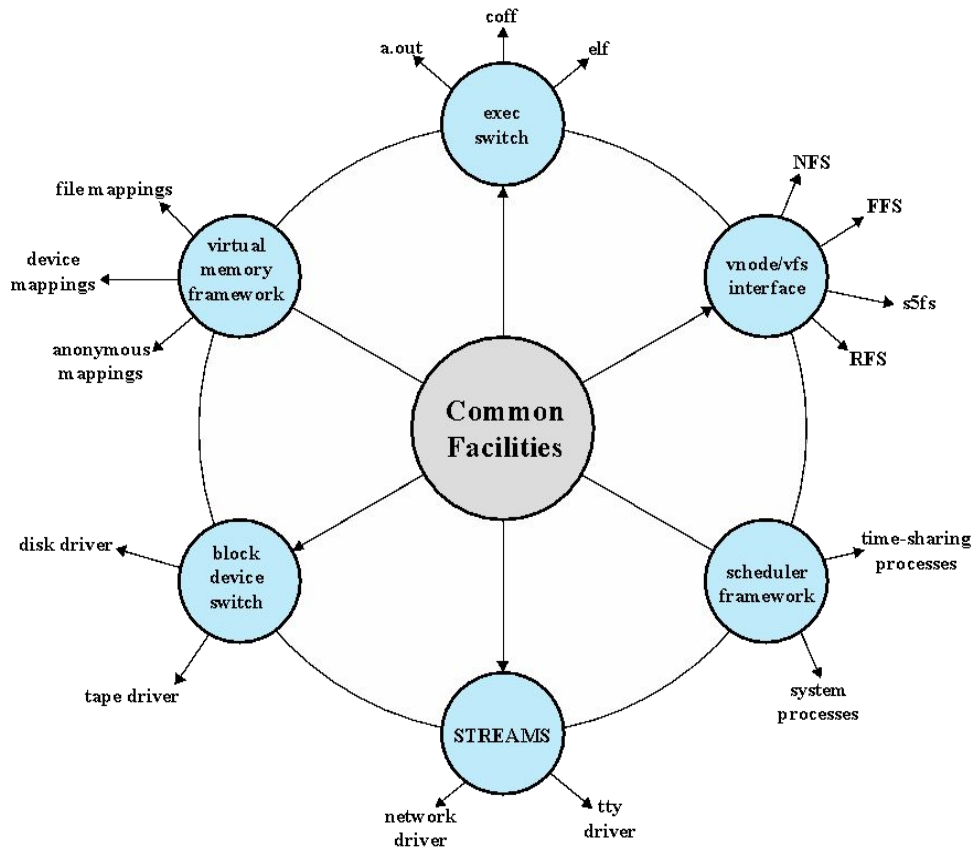
This paper discusses the nature and implementation of the file system and of the user command interface.

### I. INTRODUCTION

There have been four versions of the Unix time-sharing system.

12 The earliest (circa 1969-70) ran on the Digital Equipment Corporation PDP-7 and -9 computers. The second version ran on the unprotected PDP-11/20 computer. The third incorporated multiprogramming and ran on the PDP-11/34, /40, /45, /60, and /70 computers; it is the one described in the previously published version of this paper, and is also the most widely used today. This paper describes only the fourth, current system that runs on the PDP-11/70 and the Interdata 8/32 computers. In fact, the differences among the various systems is rather small; most of the revisions made to the originally published version of this paper, aside from those concerned with style, had to do with details of the implementation of the file system.

# Modern Unix Kernel



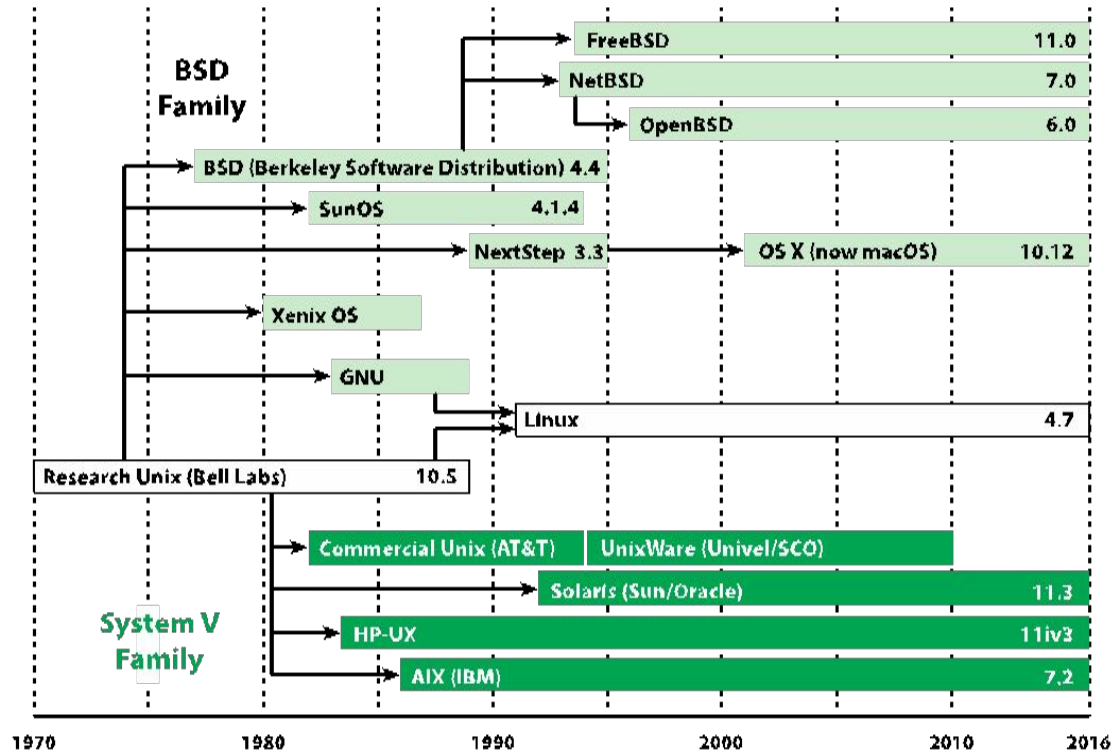



Figure 2.17 Unix Family Tree

# LINUX Overview

- Started out as a UNIX variant for the IBM PC
- Linus Torvalds, a Finnish student of computer science, wrote the initial version
- Linux was first posted on the Internet in 1991
- Today it is a full-featured UNIX system that runs on virtually all platforms
- Is free and the source code is available
- Key to the success of Linux has been the availability of free software packages under the auspices of the Free Software Foundation (FSF)
- Highly modular and easily configured



# Modular Structure – Loadable Modules

- Linux
  - does not use a microkernel approach,
  - it achieves many of the potential advantages of the approach by means of its particular modular architecture
- Linux is structured as a collection of modules, a number of which can be automatically loaded and unloaded on demand
- Relatively independent blocks
- A module is an object file whose code can be linked to and unlinked from the kernel at runtime
- A module is executed in kernel mode on behalf of the current process
- Have two important characteristics:
  - Dynamic linking
  - Stackable modules

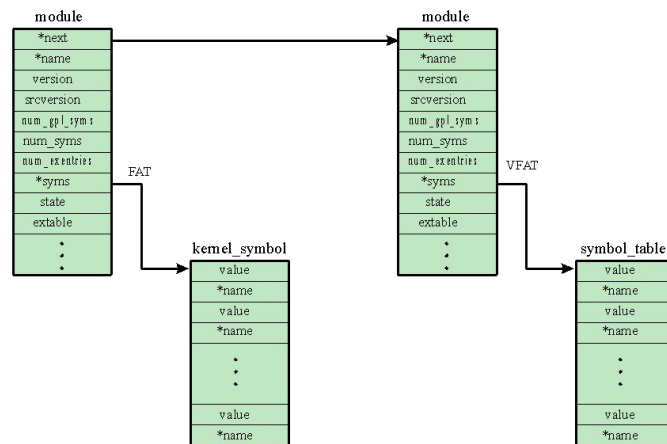


Figure 2.18 Example List of Linux Kernel Modules

# Linux kernel components

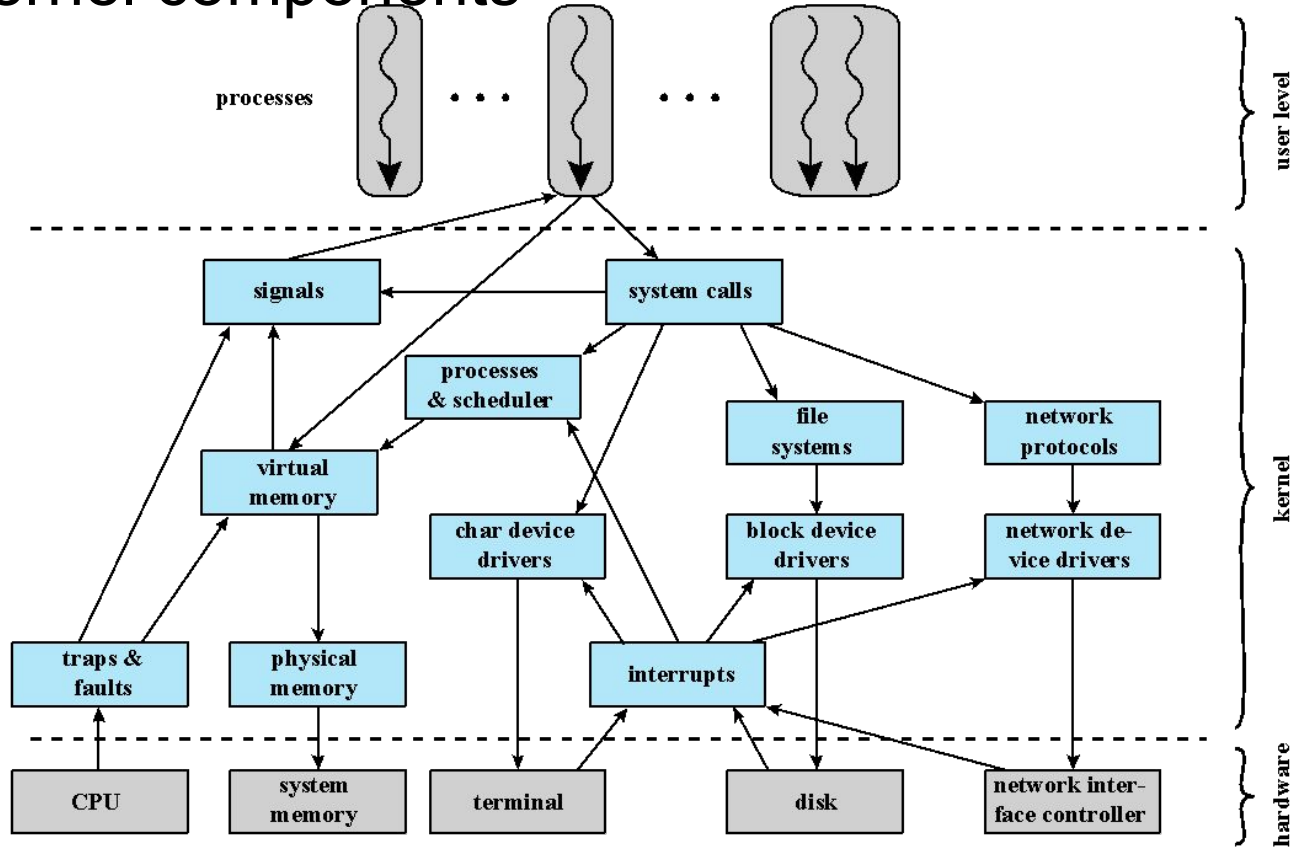


Figure 2.19 Linux Kernel Components

# Linux Signals and System Calls

SIGHUP	Terminal hangup	SIGCONT	Continue
SIGQUIT	Keyboard quit	SIGTSTP	Keyboard stop
SIGTRAP	Trace trap	SIGTTOU	Terminal write
SIGBUS	Bus error	SIGXCPU	CPU limit exceeded
SIGKILL	Kill signal	SIGVTALRM	Virtual alarm clock
SIGSEGV	Segmentation violation	SIGWINCH	Window size unchanged
SIGPIPE	Broken pipe	SIGPWR	Power failure
SIGTERM	Termination	SIGRTMIN	First real-time signal
SIGCHLD	Child status unchanged	SIGRTMAX	Last real-time signal

Filesystem related	
<b>close</b>	Close a file descriptor.
<b>link</b>	Make a new name for a file.
<b>open</b>	Open and possibly create a file or device.
<b>read</b>	Read from file descriptor.
<b>write</b>	Write to file descriptor
Process related	
<b>execve</b>	Execute program.
<b>exit</b>	Terminate the calling process.
<b>getpid</b>	Get process identification.
<b>setuid</b>	Set user identity of the current process.
<b>ptrace</b>	Provides a means by which a parent process may observe and control the execution of another process, and examine and change its core image and registers.
Scheduling related	
<b>sched_getparam</b>	Sets the scheduling parameters associated with the scheduling policy for the process identified by <code>pid</code> .
<b>sched_get_priority_max</b>	Returns the maximum priority value that can be used with the scheduling algorithm identified by <code>policy</code> .
<b>sched_setscheduler</b>	Sets both the scheduling policy (e.g., FIFO) and the associated parameters for the process <code>pid</code> .
<b>sched_rr_get_interval</b>	Writes into the <code>timespec</code> structure pointed to by the parameter <code>tp</code> the round robin time quantum for the process <code>pid</code> .
<b>sched_yield</b>	A process can relinquish the processor voluntarily without blocking via this system call. The process will then be moved to the end of the queue for its static priority and a new process gets to run.

Interprocess Communication (IPC) related	
<b>msgrcv</b>	A message buffer structure is allocated to receive a message. The system call then reads a message from the message queue specified by <code>msgid</code> into the newly created message buffer.
<b>semctl</b>	Performs the control operation specified by <code>cmd</code> on the semaphore set <code>semid</code> .
<b>semop</b>	Performs operations on selected members of the semaphore set <code>semid</code> .
<b>shmat</b>	Attaches the shared memory segment identified by <code>shmid</code> to the data segment of the calling process.
<b>shmctl</b>	Allows the user to receive information on a shared memory segment, set the owner, group, and permissions of a shared memory segment, or destroy a segment.
Socket (networking) related	
<b>bind</b>	Assigns the local IP address and port for a socket. Returns 0 for success and <code>-1</code> for error.
<b>connect</b>	Establishes a connection between the given socket and the remote socket associated with <code>sockaddr</code> .
<b>gethostname</b>	Returns local host name.
<b>send</b>	Send the bytes contained in buffer pointed to by <code>*msg</code> over the given socket.
<b>setsockopt</b>	Sets the options on a socket
Miscellaneous	
<b>fsync</b>	Copies all in-core parts of a file to disk, and waits until the device reports that all parts are on stable storage.
<b>time</b>	Returns the time in seconds since January 1, 1970.
<b>vhangup</b>	Simulates a hangup on the current terminal. This call arranges for other users to have a "clean" tty at login time.

# Summary


- Hardware types
- Operating system objectives and functions
  - User/computer interface
  - Resource manager
    - Evolution of operating systems
  - Serial processing
  - Simple/multiprogrammed/time-sharing batch systems
    - Major achievements
    - Fault tolerance
  - Fundamental concepts
  - Faults
  - OS mechanisms
- Traditional Unix systems
  - History/description
- Linux
  - History
  - Modular structure
  - Kernel components



# Process Description and Control

# Process Elements

- Resources are made available to multiple applications
  - The processor is switched among multiple applications so all will appear to be progressing
  - The processor and I/O devices can be used efficiently
- Two essential elements of a process are:
  - Code: when the processor begins to execute the program code, we refer to this executing entity as a process
  - Data: which may be shared with other processes that are executing the same program



Program code

A set of data associated with that code

Figure 3.10 Processes and Resources (resource allocation at one snapshot in time)

# Tables within the OS

- Memory tables:
  - keep track of both main (real) and secondary (virtual) memory
  - processes use some sort of virtual memory and/or simple swapping mechanism
- The memory tables must include the following information:
  - The allocation of main memory to processes
  - The allocation of secondary memory to processes
- I/O tables
  - used by the OS to manage the I/O devices and channels of the computer system.
  - At any given time, an I/O device may be available or assigned to a particular process.
  - If an I/O operation is in progress, the OS needs to know the status of the I/O operation and the location in main memory being used as the source or destination of the I/O transfer.
- File tables
  - provide information about the existence of files, their location on secondary memory, their current status, and other attributes
  - Much, if not all, of this information may be maintained by a file management system
- The OS must maintain process tables to manage processes

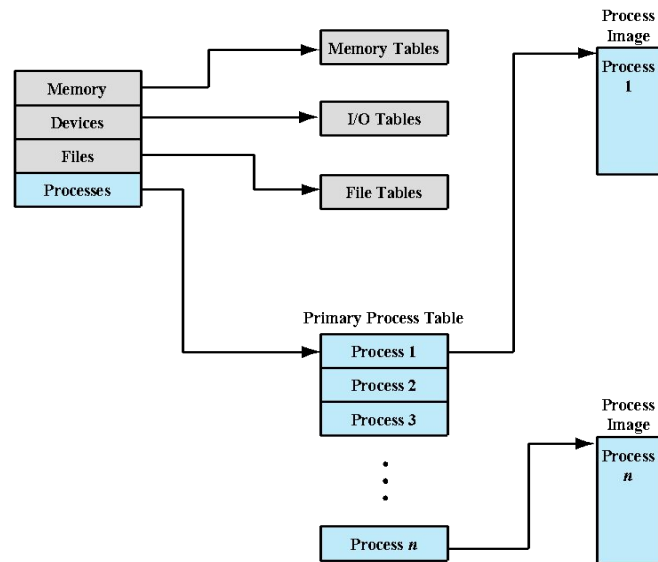
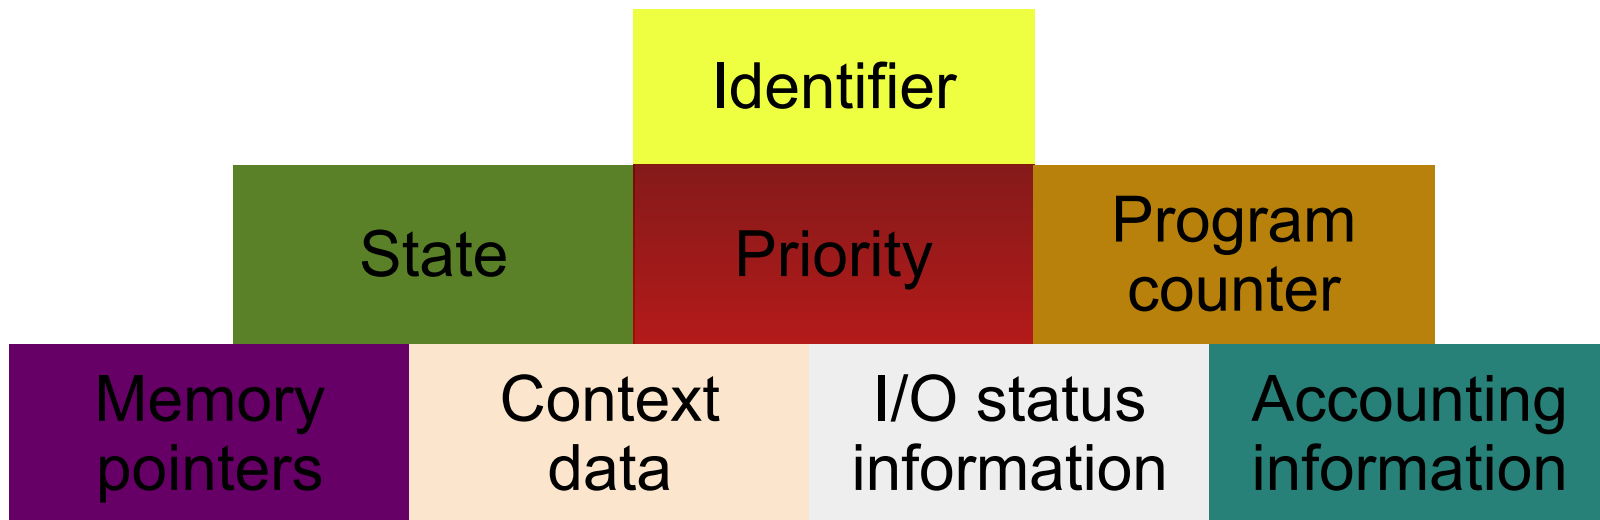


Figure 3.11 General Structure of Operating System Control Tables

# Process Elements

- While the program is executing, a process can be uniquely characterized by a number of elements, including:



# Process Control Block

- Contains the process elements
- It is possible to interrupt a running process and later resume execution as if the interruption had not occurred
- Created and managed by the operating system
- Key tool that allows support for multiple processes

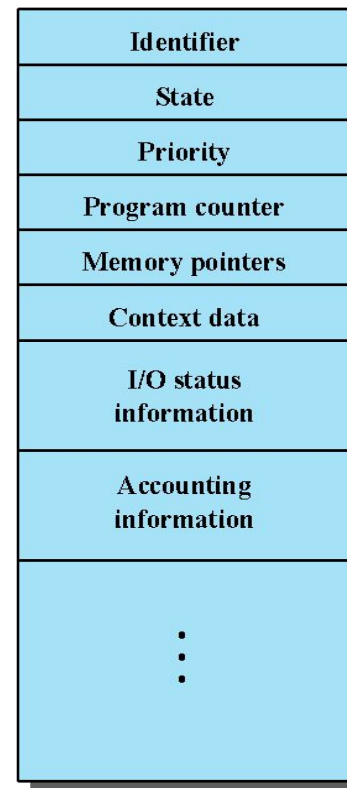


Figure 3.1 Simplified Process Control Block

# Process States

# Process States

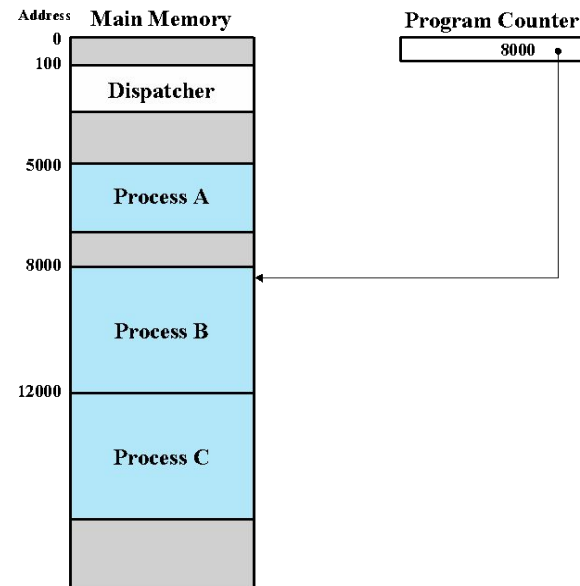
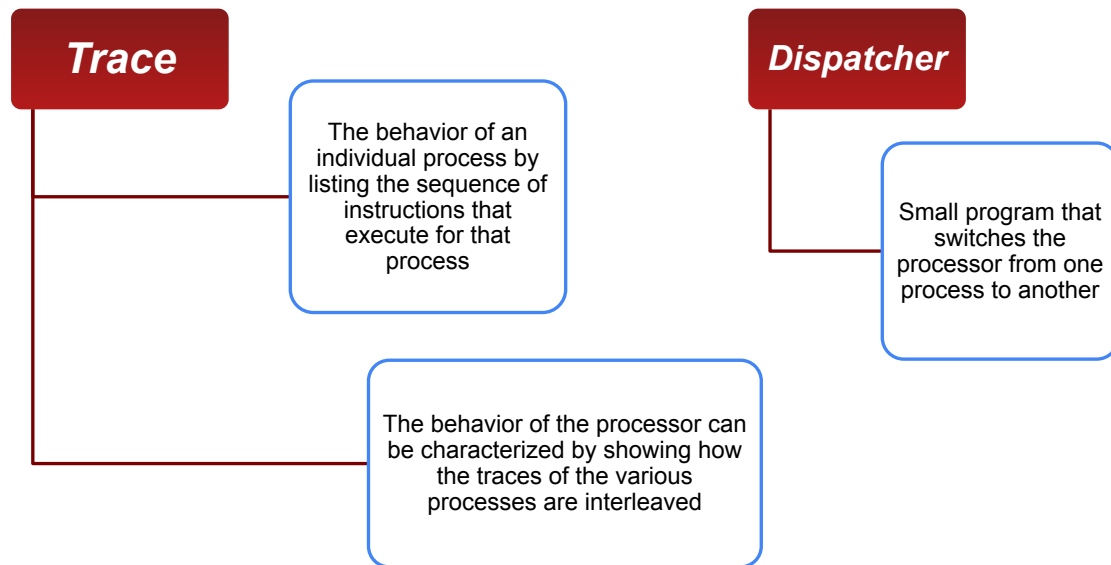
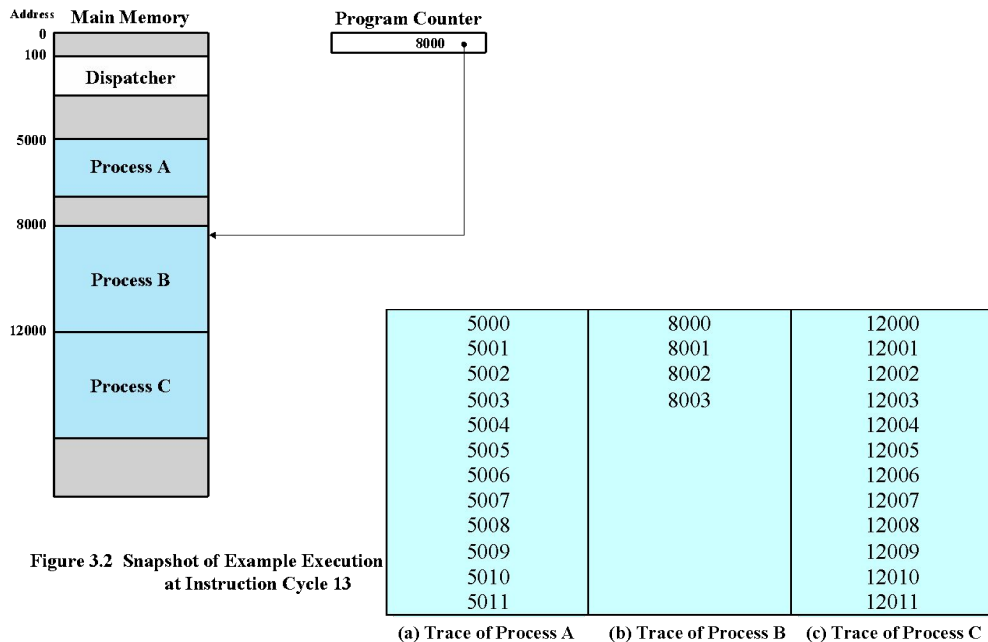


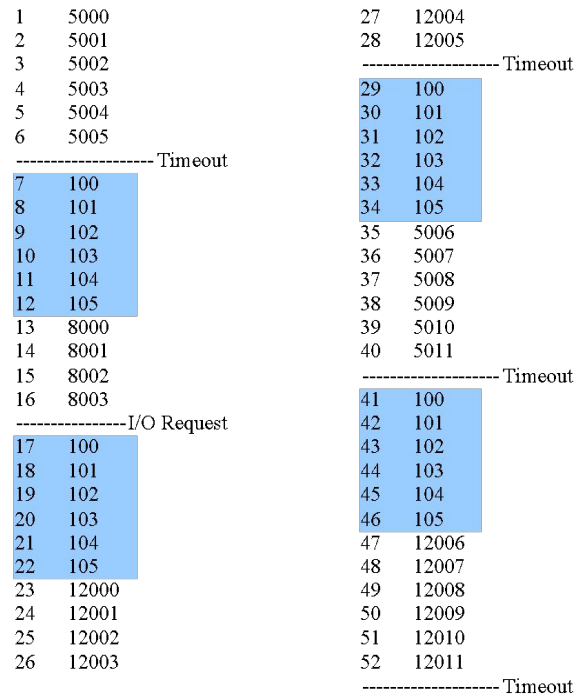
Figure 3.2 Snapshot of Example Execution (Figure 3.4) at Instruction Cycle 13

# Process Execution



5000 = Starting address of program of Process A  
 8000 = Starting address of program of Process B  
 12000 = Starting address of program of Process C

Figure 3.3 Traces of Processes of Figure 3.2




100 = Starting address of dispatcher program

Shaded areas indicate execution of dispatcher process;  
 first and third columns count instruction cycles;  
 second and fourth columns show address of instruction being executed


Figure 3.4 Combined Trace of Processes of Figure 3.2



# Two-State Process Model



(a) State transition diagram




(b) Queuing diagram

Figure 3.5 Two-State Process Model

# Process Creation

- Process spawning
  - When the OS creates a process at the explicit request of another process
  - Parent process: is the original, creating, process
  - Child process: is the new process
- Linux
  - An existing process makes an exact copy of itself
  - The child process has the same environment as its parent, only the process ID number is different. This procedure is called **forking**.
  - After the forking process, the **address space** of the child process is overwritten with the new process data. This is done through an **exec** call to the system.
  - Many programs, for instance, **daemonize** their child processes, so they can keep on running when the parent stops or is being stopped.

New batch job	The OS is provided with a batch job control stream, usually on tape or disk. When the OS is prepared to take on new work, it will read the next sequence of job control commands.
Interactive logon	A user at a terminal logs on to the system.
Created by OS to provide a service	The OS can create a process to perform a function on behalf of a user program, without the user having to wait (e.g., a process to control printing).
Spawned by existing process	For purposes of modularity or to exploit parallelism, a user program can dictate the creation of a number of processes.



# Process Termination

- There must be a means for a process to indicate its completion
- A batch job should include a HALT instruction or an explicit OS service call for termination
- For an interactive application, the action of the user will indicate when the process is completed (e.g. log off, quitting an application)
- Linux:
  - Program returns its exit status to the parent
  - This exit status is a number providing the results of the program's execution
  - The system of returning information upon executing a job has its origin in the C programming language in which UNIX has been written.

Normal completion	The process executes an OS service call to indicate that it has completed running.
Time limit exceeded	The process has run longer than the specified total time limit. There are a number of possibilities for the type of time that is measured. These include total elapsed time ("wall clock time"), amount of time spent executing, and, in the case of an interactive process, the amount of time since the user last provided any input.
Memory unavailable	The process requires more memory than the system can provide.
Bounds violation	The process tries to access a memory location that it is not allowed to access.
Protection error	The process attempts to use a resource such as a file that it is not allowed to use, or it tries to use it in an improper fashion, such as writing to a read-only file.
Arithmetic error	The process tries a prohibited computation, such as division by zero, or tries to store numbers larger than the hardware can accommodate.
Time overrun	The process has waited longer than a specified maximum for a certain event to occur.
I/O failure	An error occurs during input or output, such as inability to find a file, failure to read or write after a specified maximum number of tries (when, for example, a defective area is encountered on a tape), or invalid operation (such as reading from the line printer).
Invalid instruction	The process attempts to execute a nonexistent instruction (often a result of branching into a data area and attempting to execute the data).
Privileged instruction	The process attempts to use an instruction reserved for the operating system.
Data misuse	A piece of data is of the wrong type or is not initialized.
Operator or OS intervention	For some reason, the operator or the operating system has terminated the process (e.g., if a deadlock exists).
Parent termination	When a parent terminates, the operating system may automatically terminate all of the offspring of that parent.
Parent request	A parent process typically has the authority to terminate any of its offspring.

# Five-State Process Model

- **New:** A process that has just been created but has not yet been admitted to the pool of executable processes by the OS. Typically, a new process has not yet been loaded into main memory, although its process control block has been created.
- **Ready:** A process that is prepared to execute when given the opportunity.
- **Running:** The process that is currently being executed.
- **Blocked/Waiting:** A process that cannot execute until some event occurs, such as the completion of an I/O operation.
- **Exit:** A process that has been released from the pool of executable processes by the OS, either because it halted or because it aborted for some reason.




Figure 3.6 Five-State Process Model

# Example: Trace of Process Execution

1	5000	27	12004
2	5001	28	12005
3	5002		-----Timeout
4	5003	29	100
5	5004	30	101
6	5005	31	102
	-----Timeout	32	103
7	100	33	104
8	101	34	105
9	102	35	5006
10	103	36	5007
11	104	37	5008
12	105	38	5009
13	8000	39	5010
14	8001	40	5011
15	8002		-----Timeout
16	8003	41	100
	-----I/O Request	42	101
17	100	43	102
18	101	44	103
19	102	45	104
20	103	46	105
21	104	47	12006
22	105	48	12007
23	12000	49	12008
24	12001	50	12009
25	12002	51	12010
26	12003	52	12011
	-----Timeout		




Figure 3.7 Process States for Trace of Figure 3.4

100 = Starting address of dispatcher program

Shaded areas indicate execution of dispatcher process;

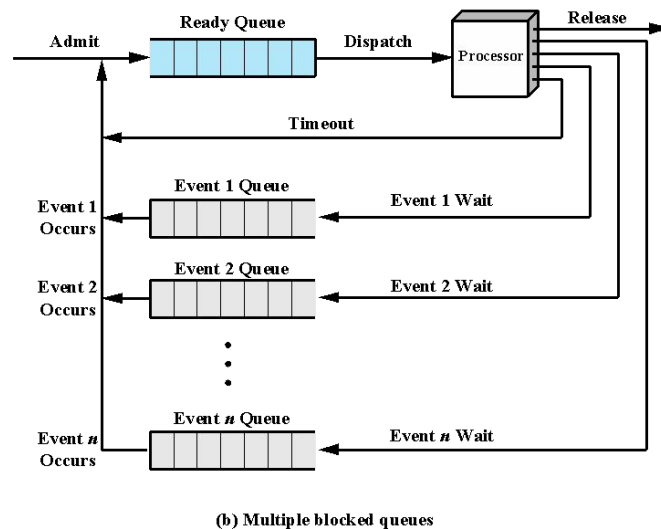

first and third columns count instruction cycles;

second and fourth columns show address of instruction being executed

Figure 3.4 Combined Trace of Processes of Figure 3.2


# Faster/Efficient Scheduling of Processes

- Example, Linux Scheduler
  - separate run queue for each processor
  - each processor only selects processes from its own queue to run
  - it's possible for one processor to be idle while others have jobs waiting in their run queues
  - Periodically, the queues are rebalanced:
    - if one processor's run queue is too long
    - some processes are moved from it to another processor's queue
- 140 separate queues, one for each priority level
  - Number can be changed at a given site
  - Two sets, active and expired
  - Priorities 0-99 for real-time processes
  - Priorities 100-139 for normal processes

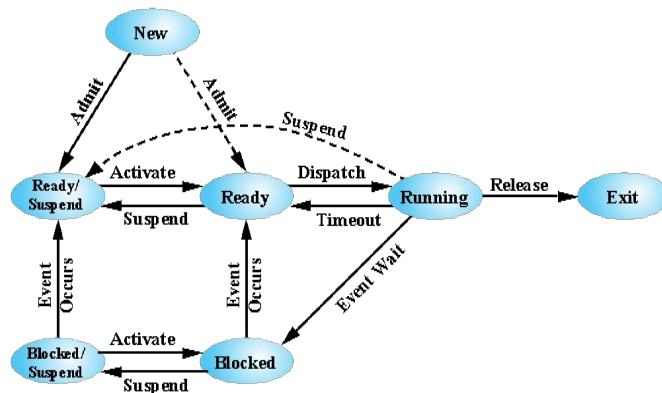


# Suspended Processes - Swapping

- Swapping
  - is an I/O operation and therefore there is the potential for making the problem worse, not better
  - disk I/O is generally the fastest I/O on a system
  - swapping will usually enhance performance.
- Moving part of all of a process from main memory to disk
  - When none of the processes in main memory is in the Ready state, the OS swaps one of the blocked processes out on to disk into a suspend queue
  - This is a queue of existing processes that have been temporarily kicked out of main memory, or suspended
  - The OS then brings in another process from the suspend queue or it honors a new-process request
  - Execution then continues with the newly arrived process



(a) With One Suspend State



(b) With Two Suspend States

Figure 3.9 Process State Transition Diagram with Suspend States

# Characteristics of a Suspended Process

- The process may or may not be waiting on an event
- The process is not immediately available for execution
- The process was placed in a suspended state by an agent: either itself, a parent process, or the OS, for the purpose of preventing its execution
- The process may not be removed from this state until the agent explicitly orders the removal

Swapping	The OS needs to release sufficient main memory to bring in a process that is ready to execute.
Other OS reason	The OS may suspend a background or utility process or a process that is suspected of causing a problem.
Interactive user request	A user may wish to suspend execution of a program for purposes of debugging or in connection with the use of a resource.
Timing	A process may be executed periodically (e.g., an accounting or system monitoring process) and may be suspended while waiting for the next time interval.
Parent process request	A parent process may wish to suspend execution of a descendent to examine or modify the suspended process, or to coordinate the activity of various descendants.



# Process Control

# Process Control Structures

- A process must include a program or set of programs to be executed
  - A process will consist of at least sufficient memory to hold the programs and data of that process
  - The execution of a program typically involves a stack that is used to keep track of procedure calls and parameter passing between procedures
- Each process has associated with it a number of attributes that are used by the OS for process control
  - The collection of program, data, stack, and attributes is referred to as the process image
  - Process image location will depend on the memory management scheme being used

To manage and control a process the OS must know:

- Where the process is located
- The attributes of the process that are necessary for its management

# Typical Elements of a Process Image

## User Data

The modifiable part of the user space. May include program data, a user stack area, and programs that may be modified.

## User Program

The program to be executed.

## Stack

Each process has one or more last-in-first-out (LIFO) stacks associated with it. A stack is used to store parameters and calling addresses for procedure and system calls.

## Process Control Block

Data needed by the OS to control the process (see Table 3.5).

## Process Control Information

### Scheduling and State Information

This is information that is needed by the operating system to perform its scheduling function. Typical items of information:

- Process state:** Defines the readiness of the process to be scheduled for execution (e.g., running, ready, waiting, halted).
- Priority:** One or more fields may be used to describe the scheduling priority of the process. In some systems, several values are required (e.g., default, current, highest-allowable)
- Scheduling-related information:** This will depend on the scheduling algorithm used. Examples are the amount of time that the process has been waiting and the amount of time that the process executed the last time it was running.
- Event:** Identity of event the process is awaiting before it can be resumed.

### Data Structuring

A process may be linked to other process in a queue, ring, or some other structure. For example, all processes in a waiting state for a particular priority level may be linked in a queue. A process may exhibit a parent-child (creator-created) relationship with another process. The process control block may contain pointers to other processes to support these structures.

### Interprocess Communication

Various flags, signals, and messages may be associated with communication between two independent processes. Some or all of this information may be maintained in the process control block.

### Process Privileges

Processes are granted privileges in terms of the memory that may be accessed and the types of instructions that may be executed. In addition, privileges may apply to the use of system utilities and services.

### Memory Management

This section may include pointers to segment and/or page tables that describe the virtual memory assigned to this process.

### Resource Ownership and Utilization

Resources controlled by the process may be indicated, such as opened files. A history of utilization of the processor or other resources may also be included; this information may be needed by the scheduler.

# Process Identification

- Each process is assigned a unique numeric identifier
  - Otherwise there must be a mapping that allows the OS to locate the appropriate tables based on the process identifier
- Tables controlled by the OS may use process identifiers to cross-reference process tables
- Memory tables may be organized to provide a map of main memory with an indication of which process is assigned to each region
  - Similar references will appear in I/O and file tables
  - When processes communicate with one another, the process identifier informs the OS of the destination of a particular communication
- When processes are allowed to create other processes, identifiers indicate the parent and descendents of each process

## Process Identification

### Identifiers

Numeric identifiers that may be stored with the process control block include

- Identifier of this process
- Identifier of the process that created this process (parent process)
- User identifier

## Processor State Information

### User-Visible Registers

A user-visible register is one that may be referenced by means of the machine language that the processor executes while in user mode. Typically, there are from 8 to 32 of these registers, although some RISC implementations have over 100.

### Control and Status Registers

These are a variety of processor registers that are employed to control the operation of the processor. These include

- **Program counter:** Contains the address of the next instruction to be fetched
- **Condition codes:** Result of the most recent arithmetic or logical operation (e.g., sign, zero, carry, equal, overflow)
- **Status information:** Includes interrupt enabled/disabled flags, execution mode

### Stack Pointers

Each process has one or more last-in-first-out (LIFO) system stacks associated with it. A stack is used to store parameters and calling addresses for procedure and system calls. The stack pointer points to the top of the stack.